

# Multiclass Classifiers

One vs All and All Pairs

Angela Zhu, Michael Lu, Qiming Fang, Jessica Wan

# Math

- Representation

$$h_w(x)_j = \frac{e^{\langle w_j, x \rangle}}{\sum_{s=1}^k e^{\langle w_s, x \rangle}}$$

- Loss

$$L_s(h_w) = -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^k \mathbf{1}[y_i = j] \log h_w(x_i)_j$$

- Optimizer

- Stochastic gradient descent

$$w = w - \alpha \nabla L_s(h_w)$$

# Stochastic Gradient Descent

input:

training examples  $X, Y$  of size  $S$

step size  $\alpha$

batch size  $b < S$

converge = False

while not converge:

    shuffle  $X, Y$

    prev\_epoch\_loss = loss( $X, Y$ )

    for each batch:

$L_w = \text{zeros}(\text{num\_classes}, \text{num\_features} + 1)$

        for  $x, y$  in batch:

            probabilities = softmax( $x$ )

            for class  $j$  in classes:

$h_{wx} = \text{probabilities}[j]$

                if  $y == j$ :  $L_w[j] += (h_{wx} - 1) * x$

                else:  $L_w[j] += h_{wx} * x$

$w -= \alpha * L_w / \text{num\_examples\_in\_batch}$

    curr\_epoch\_loss = loss( $X, Y$ )

    converge =  $\text{diff}(\text{curr\_epoch\_loss}, \text{prev\_epoch\_loss}) < \text{threshold}$

# Algorithms - Pseudo Code

## All Pairs

- Train one binary classifier for each class pair.
- Use the argmax of the average of the binary classifiers to predict.

```
input:
  training set  $S = (x_1, y_1), \dots, (x_m, y_m)$ 
  algorithm for binary classification  $A$ 
foreach  $i, j \in \mathcal{Y}, i < j$ :
   $S_{i,j}$  initialize to be empty
  for  $t=1, \dots, m$ :
    if  $y_t = i$ , add  $(x_t, 1)$  to  $S_{i,j}$ 
    if  $y_t = j$ , add  $(x_t, -1)$  to  $S_{i,j}$ 
  let  $h_{i,j} = A(S_{i,j})$ 
output:
  the multi-class hypothesis defined by  $h(x) \in \operatorname{argmax}_{i \in \mathcal{Y}} (\sum_{j \in \mathcal{Y}} \operatorname{sign}(j - i) h_{i,j}(x))$ 
```

## One-vs-All

- Train one binary classifier for each class which predicts whether or not a sample belongs to that class (1) or not (0).
- Select the result of the classifier with the largest positive prediction.

```
input:
  training set  $S = (x_1, y_1), \dots, (x_m, y_m)$ 
  algorithm for binary classification model
foreach  $i \in \mathcal{Y}$ 
  let  $S_i = ((x_1, (1)^{\mathbb{1}[y_1 \neq i]}), \dots, (x_m, (1)^{\mathbb{1}[y_m \neq i]}))$ 
  let  $h_i = A(S_i)$ 
output:
  the multi-class hypothesis defined by  $h(x) \in \operatorname{argmax}_{i \in \mathcal{Y}} h_i(x)$ 
```

# Previous Work

## Iris Data + German Numerical Credit Data



- Combinations of our model + our LR, our model + SK LR, and SK Model + SK LR
- Iris Data:
  - 120 training size, 30 testing size
  - 4 features + 1 bias
  - 3 classes
- Credit Data:
  - 350 training size, 150 testing size
  - 69 features + 1 bias
  - 2 classes

# Previous Work

## Iris Data + German Numerical Credit Data



**Iris Data**

	All Pairs	One V All
SKLearn	<b>0.9</b>	<b>0.833</b>
Our Model	<b>0.933</b>	<b>0.966</b>

**Credit Data**

	All Pairs	One V All
SKLearn	<b>0.7028</b>	<b>0.7028</b>
Our Model	<b>0.7028</b>	<b>0.7028</b>

# Summary

- Checked that our One vs All and All Pairs algorithms were correctly implemented no matter what our underlying binary classification algorithm A was
  - German credit data: had two classes
    - Our own One vs All and All Pairs had same predictions and accuracy with our own LogisticRegression binary classification algo, and with sklearn's binary classification LogisticRegression algo
  - Unit tests:
    - Used sklearn's OneVsRestClassifier and OneVsOneClassifier with underlying binary classification sklearn's LogisticRegression algo against our own multiclass classifiers with sklearn's LogisticRegression as the underlying binary classification
- Our own LogisticRegression and sklearn's LogisticRegression algo are implemented differently, however:
  - Iris data: both our models One vs All and All Pairs with our own LogisticRegression and sklearn's OneVsRest and OneVsOne classifiers with sklearn's LogisticRegression give similar accuracies across multiple runs

# Reflection

- Challenges:
  - Checking to see if our LogisticRegression for binary classification is accurate
    - Sklearn's LogisticRegression default uses L-BFGS optimization algorithm to find minimum of a function, other optimization algorithms available would be Stochastic Average Gradient, which would still be different from our implementation of Stochastic Gradient Descent
  - Debugging: for two classes One vs All and All Pairs should be the same
- Interesting to see how different optimization algorithms for the underlying binary classification algorithm would affect accuracy levels for One vs All and All Pairs multiclass algorithms
  - On Iris dataset, if we train and test on a random 80/20 portion of the dataset from shuffling, across multiple runs sometimes SGD is better sometimes L-BFGS is better